

RSS Collector

An offline-friendly Google News RSS ingestion pipeline designed for LLM integration. This tool fetches Google News RSS feeds based on configurable keywords, deduplicates articles intelligently, and stores them in a structured format suitable for downstream processing.

Features

- **Keyword-based RSS fetching** from Google News
- **Intelligent deduplication** based on URL and content hashing
- **Scheduled execution** with configurable timing
- **Offline compatibility** with single-port external communication
- **Robust error handling** with retry mechanisms
- **LLM-ready JSON output** with clean, structured data
- **Comprehensive logging** for monitoring and debugging

Requirements

- Python 3.10 or higher
- RHEL-compatible environment (or any Linux distribution)
- Internet access through a single configurable port
- Optional: Proxy server (Tinyproxy or Squid)

Installation

Option 1: Using pip

```
bash

# Clone the repository
git clone <repository-url>
cd rss_collector

# Install the package
pip install -e .

# For development with testing dependencies
pip install -e ".[dev]"
```

Option 2: Manual setup

```
bash
```

```
# Install dependencies
```

```
pip install -r requirements.txt
```

```
# Create necessary directories
```

```
mkdir -p feeds logs config
```

```
# Copy configuration files
```

```
cp config/feeds.json.example config/feeds.json
```

```
cp config/settings.json.example config/settings.json
```

Configuration

1. Keywords Configuration (config/feeds.json)

```
json
```

```
{
  "keywords": [
    "artificial intelligence",
    "python programming",
    "technology news"
  ],
  "keyword_groups": [
    {
      "name": "Programming",
      "keywords": ["python", "javascript", "software development"]
    },
    {
      "name": "AI/ML",
      "keywords": ["machine learning", "deep learning"]
    }
  ]
}
```

2. System Settings (config/settings.json)

json

```
{
  "http": {
    "timeout": 30,
    "max_retries": 3,
    "retry_delay": 5,
    "user_agent": "RSS-Collector/1.0"
  },
  "proxy": {
    "enabled": true,
    "host": "localhost",
    "port": 8081,
    "protocol": "http"
  },
  "schedule": [
    {
      "time": "05:00",
      "description": "Morning news fetch"
    },
    {
      "time": "14:00",
      "description": "Afternoon news fetch"
    }
  ],
  "storage": {
    "base_dir": "./",
    "feeds_dir": "feeds",
    "logs_dir": "logs",
    "cleanup_days": 30
  },
  "processing": {
    "group_delay_minutes": 5,
    "max_articles_per_feed": 100
  }
}
```

3. Environment Variables (.env)

```
bash
```

```
# Proxy configuration
```

```
HTTP_PROXY=http://localhost:8081
```

```
HTTPS_PROXY=http://localhost:8081
```

```
# Logging level
```

```
LOG_LEVEL=INFO
```

```
# Override settings
```

```
RSS_BASE_DIR=/custom/path
```

```
RSS_MAX_RETRIES=5
```



Usage

Running Once

```
bash
```

```
# Run the collector once
```

```
python run.py
```

```
# Or using the installed command
```

```
rss-collector-run
```

Scheduled Execution

```
bash
```

```
# Start the scheduler (runs continuously)
```

```
python -m src.main --scheduled
```

```
# Or using the entry point
```

```
rss-collector --scheduled
```

Command Line Options

```
bash
```

```
# Show help
```

```
python run.py --help
```

```
# Run with custom config
```

```
python run.py --config /path/to/config
```

```
# Run specific keywords only
```

```
python run.py --keywords "python,ai,technology"
```

```
# Verbose logging
```

```
python run.py --verbose
```

```
# Dry run (no actual fetching)
```

```
python run.py --dry-run
```

Output Structure

The collector creates the following directory structure:

```
rss_collector/
├── feeds/
│   ├── 2024-01-15.json      # Daily feeds
│   ├── 2024-01-16.json
│   └── ...
├── logs/
│   ├── 2024-01-15.log      # Daily logs
│   ├── 2024-01-16.log
│   └── ...
└── config/
    ├── feeds.json          # Keywords config
    └── settings.json       # System settings
```

JSON Output Format

json

```
[
  {
    "fetched_at": "2024-01-15T08:00:00Z",
    "query": "artificial intelligence",
    "source_url": "https://news.google.com/rss/search?q=artificial%20intelligence&hl=en",
    "articles": [
      {
        "title": "AI Breakthrough in Natural Language Processing",
        "link": "https://example.com/ai-breakthrough",
        "published": "Mon, 15 Jan 2024 07:30:00 GMT",
        "source": "TechNews",
        "snippet": "Researchers achieve new milestone in AI language understanding..."
      }
    ]
  }
]
```

Proxy Setup

Using Tinyproxy

```
bash

# Install tinyproxy
sudo yum install tinyproxy

# Configure /etc/tinyproxy/tinyproxy.conf
Port 8081
Listen 127.0.0.1
Allow 127.0.0.1

# Start service
sudo systemctl start tinyproxy
sudo systemctl enable tinyproxy
```

Using iptables (Alternative)

```
bash

# Route traffic through specific port
sudo iptables -t nat -A OUTPUT -p tcp --dport 80 -j REDIRECT --to-port 8081
sudo iptables -t nat -A OUTPUT -p tcp --dport 443 -j REDIRECT --to-port 8081
```

Testing

bash

Run all tests

```
python -m pytest
```

Run with coverage

```
python -m pytest --cov=src
```

Run specific test file

```
python -m pytest tests/test_rss_fetcher.py
```

Run with verbose output

```
python -m pytest -v
```

Logging

The application provides comprehensive logging:

- **INFO:** Normal operations, fetch statistics
- **WARNING:** Recovered errors, configuration issues
- **ERROR:** Failed requests, parsing errors
- **DEBUG:** Detailed execution information

Log files are created daily in the `logs/` directory with automatic rotation.

Monitoring

Check Status

```
bash
```

```
# View recent logs
```

```
tail -f logs/$(date +%Y-%m-%d).log
```

```
# Check feed files
```

```
ls -la feeds/
```

```
# View statistics
```

```
python -c "
```

```
import json
```

```
from datetime import datetime
```

```
date = datetime.now().strftime('%Y-%m-%d')
```

```
with open(f'feeds/{date}.json') as f:
```

```
    data = json.load(f)
```

```
    print(f'Total feeds: {len(data)}')
```

```
    print(f'Total articles: {sum(len(feed["articles"]) for feed in data)}')
```

```
"
```

Health Check Script

```
bash
```

```
#!/bin/bash
```

```
# health_check.sh
```

```
LOG_FILE="logs/$(date +%Y-%m-%d).log"
```

```
FEED_FILE="feeds/$(date +%Y-%m-%d).json"
```

```
if [ -f "$LOG_FILE" ] && [ -f "$FEED_FILE" ]; then
```

```
    echo "RSS Collector: OK"
```

```
    echo "Last run: $(tail -1 $LOG_FILE | cut -d' ' -f1-2)"
```

```
    echo "Articles today: $(jq '[.[].articles | length] | add' $FEED_FILE)"
```

```
else
```

```
    echo "RSS Collector: ERROR - Missing files"
```

```
    exit 1
```

```
fi
```

Troubleshooting

Common Issues

1. Connection Errors

- Check proxy configuration
- Verify firewall settings
- Test internet connectivity

2. Parse Errors

- Check RSS feed validity
- Verify feedparser version
- Review error logs

3. Storage Issues

- Check disk space
- Verify directory permissions
- Review file system errors

4. Scheduling Problems

- Check system timezone
- Verify cron permissions
- Review scheduler logs

Debug Mode

bash

```
# Enable debug logging
```

```
export LOG_LEVEL=DEBUG
```

```
python run.py --verbose
```

```
# Check configuration
```

```
python -c "
```

```
from src.config_manager import ConfigManager
```

```
config = ConfigManager()
```

```
print('Config loaded successfully')
```

```
print(f'Keywords: {len(config.keywords)}')
```

```
print(f'Schedule: {len(config.schedule)}')
```

```
"
```

Performance

Typical Performance Metrics

- **Fetch time:** 2-5 seconds per keyword
- **Parse time:** <1 second per RSS feed
- **Memory usage:** 50-100 MB during execution
- **Storage:** ~1-2 MB per 100 articles

Optimization Tips

1. Adjust `group_delay_minutes` for rate limiting

2. Set `max_articles_per_feed` to control memory usage
3. Use `cleanup_days` to manage disk space
4. Monitor proxy performance for bottlenecks

Contributing

1. Fork the repository
2. Create a feature branch
3. Make your changes
4. Add tests for new functionality
5. Run the test suite
6. Submit a pull request

Development Setup

```
bash

# Clone and setup development environment
git clone <repository-url>
cd rss_collector
pip install -e ".[dev]"

# Pre-commit hooks
pre-commit install

# Run tests
pytest

# Code formatting
black src/ tests/
flake8 src/ tests/
```

License

This project is licensed under the MIT License. See the [LICENSE](#) file for details.

Support

For issues and questions:

1. Check the [troubleshooting section](#)
2. Review [existing issues](#)
3. Create a new issue with detailed information

Changelog

Version 1.0.0

- Initial release
 - Google News RSS fetching
 - Intelligent deduplication
 - Scheduled execution
 - Comprehensive logging
 - LLM-ready output format
-

Note: This tool is designed for educational and research purposes. Ensure compliance with Google's Terms of Service and robots.txt when using this collector.